

Evaluation Set & Scorecard — Retail / E-commerce

Author: Rishi — Intern 3 (Retail) **Purpose:** 25–30 hand-written questions with expected behaviour, + the scorecard the harness produces. Baseline-gate deliverable — exists **before** the chatbot does.

[TODO] = fill once the real corpus + dataset exist (expected doc titles, real customer IDs). Questions are written now; expected sources are pinned when the corpus is built.

A. Corpus plan (phase 1) — source your own, 15–40 docs

- **Majority real public documentation.** Retailers publish extensive help centres, returns/refund/warranty/delivery terms. Public pages only. Sources listed in `data/sources.md`.
- **Four companies — Amazon, Best Buy, IKEA, Target** — chosen for differing terminology for the same thing + genuinely conflicting policies, the mess the system must handle. Source URLs in `data/sources.md`.
- **Deliberately contradicting pair: Amazon ~30-day standard returns vs Best Buy 15-day (14-day cellular; \$45 restocking fee on activatable devices).** Bonus spread: IKEA 365-day, Target 90-day. The system must surface the conflict, not blend them. All 22 corpus docs are **real public pages** (no LLM-generated filler); their natural inconsistency in terms and coverage is the mess the system must handle.
- **Mix long and short docs** — uniform length hides chunking problems.

B. Dataset plan (phase 2) — generate it, synthetic

- ~4–6 SQLite tables: `orders`, `line_items`, `shipments`, `returns`, `customers`. A few hundred to a few thousand rows.

- **Consistent with the documents** — if a policy doc says refunds take 5 working days, the data shows ~5, with a few exceptions (exceptions are what agents call about).
 - **Awkward cases included:** a duplicate charge, a partially shipped order, a return already under review, a customer with two accounts.
 - **3–4 reference customers** documented with their IDs for demos. **[TODO: pin IDs once generated.]**
 - Small enough to inspect by hand — I must know the right answer to score my own eval.
-

C. Questions

Phrased the way an agent asks mid-call. Each records what *should* happen.

`expected_source` is **[TODO]** until the corpus is built.

DOCUMENT (7) — INCL. TWO-DOCUMENT ANSWERS + VOCABULARY DIFFERING FROM SOURCE

#	Question (agent voice)	Should do	Expected
1	"how long till they get their money back on a return?"	search	refund-window passage; note "money back" must match a "refund" chunk (vocab differs)
2	"customer's parcel never showed up — what's our process?"	search	delivery-failure passage
3	"can they send back opened electronics?"	search	returns-eligibility passage
4	"what's covered under warranty and for how long?"	search	warranty passage
5	"they want to return after 40 days, are we allowed?"	search → refuse/deny	return-window passage; answer = no
6	"returns window AND who pays return shipping?"	search (two documents)	returns-policy + shipping-cost docs
7	"does every store charge a 'restocking fee' on returns?"	search (vocab conflict across docs)	Best Buy \$45 activatable-device fee vs IKEA / Target (no such fee) — flag they differ

RECORD (6) — EXACT LOOKUPS, FILTERED LISTS, ≥1 AGGREGATE

#	Question	Should do	Expected
8	"status of order [REF]?"	query_orders	that order's status
9	"did order [REF] actually ship?"	query_shipments	shipment row
10	"list this customer's open returns"	query_returns (filtered)	filtered list
11	"how many orders has [customer] placed this year?"	query (aggregate/count)	correct count
12	"how much has this customer been refunded this year?"	query (aggregate/sum)	correct total
13	"which line items in order [REF] were delivered?"	query_line_items (filtered)	filtered rows

COMPOSITE (4) — NEED DOCUMENTS AND RECORDS IN ONE ANSWER

#	Question	Should do	Expected
14	"I was charged twice — is that allowed and did it actually happen?"	search (payments policy) + query (duplicate charge)	policy + the duplicate-charge row
15	"can they return order [REF] — what's the window and is it eligible?"	query_orders + search (returns policy)	order date + policy → eligible/not
16	"parcel split into two — is partial delivery covered, and what shipped?"	search (delivery policy) + query_shipments	policy + partial-shipment rows
17	"refund on [REF] — how long should it take and did it go through?"	search + query_returns	policy window + return status

CROSS-SERVER (4) — ≥ 2 NEED ANOTHER INDUSTRY, ≥ 1 COMPARATIVE ACROSS TWO AT ONCE

#	Question	Should do	Expected
18	"does the bank show a refund for this charge yet?"	route → Banking server	banking record; not retail
19	"customer disputing a telecom bill, not our order"	route → Telecom server	telecom; retail refuses its part
20	"do refund timelines differ between us and the bank?"	comparative — retail + Banking at once	both policies, compared
21	"compare our return window with the hotel's cancellation window"	comparative — retail + Hospitality	both, compared

ACTION (4) — INCL. ONE MISSING-FIELD → CLIENT MUST ASK

#	Question	Should do	Expected
22	"open a return for order [REF], item [ITEM], reason damaged"	<code>kb_retail_create_return</code> , confirm first	confirmation showing fields, then RMA ref
23	"start a return for this customer" (missing order/item)	client asks for the missing fields , does not invent	prompts for order_id + line_item_id
24	"raise the return we just discussed" (follow-up, multi-turn)	<code>kb_retail_create_return</code> using prior turn's order, confirm	RMA ref
25	"open a return on order [REF]" but item already returned	<code>kb_retail_create_return</code> → error (<code>retryable: false</code>)	loud fail, no silent default

UNANSWERABLE (3) — RETRIEVAL SHOULD FAIL, SYSTEM SHOULD REFUSE

#	Question	Should do	Expected
26	"what's the CEO's mobile number?"	refuse	not in data
27	"will this product be cheaper next month?"	refuse	can't predict
28	"status of order 99999999?" (no such order)	query → empty → refuse honestly	zero rows, not invented

(28 questions — inside 25–30.)

D. Scorecard — reported TWICE (baseline when retrieval works, final after tuning)

Each layer measured separately.

Retrieval (phase 1): Recall@5 $\geq 85\%$ · Recall@1 $\geq 60\%$ (*may argue accept-set on Recall@1 if my corpus has genuine policy conflicts — see design doc §7*)

Routing: correct server $\geq 90\%$ · correct tool-type $\geq 90\%$ · spurious calls $\leq 1/\text{query avg}$ · cross-server synthesis $\geq 80\%$ · composite handling $\geq 80\%$

Structured (phase 2): query correctness $\geq 90\%$ · numerical accuracy 100% · empty-result honesty 100%

Answer quality: groundedness 100% · citation accuracy 100% · correct refusal 100% · false refusal $\leq 10\%$

Action safety (phase 3): spurious writes 0 · fabricated fields 0 · correct action routing 100% · confirmation shown 100%

Latency: e2e p50 $\leq 4\text{s}$ · p95 $\leq 10\text{s}$ · retrieval $\leq 300\text{ms}$ · query $\leq 100\text{ms}$ · MCP overhead $\leq 100\text{ms}$ · (*warm vs cold reported separately — Ollama unloads idle models*)

Token efficiency: tokens/query reported → reduced $\geq 40\%$ vs naive raw-JSON injection
 · no groundedness/recall regression

Robustness (pass/fail, none may crash): 1 server down · all servers down · empty query · 5000-word query · wrong-language query · off-topic question · 10000-row match
 · missing customer ref · two identical concurrent requests

E. Measurement harness (`eval/run_evals.py`)

One script, one command, prints the scorecard as one table.

- Reads this eval set (question + expected).
- Runs **each question in a fresh session** (else a question passes only because an earlier turn retrieved the right passage → numbers stop being repeatable).
- Scores each layer separately; logs per query which servers/tools were called, what they returned, latency per stage.
- **Instruments token counts from the baseline run onward** — a reduction figure with no starting point is not a measurement.
- Runs the token test **both ways** (verbose vs compact) and reports both numbers.

(Harness code lands at the baseline gate, with the first scorecard.)